

# GPR Kernels

Mark Blyth

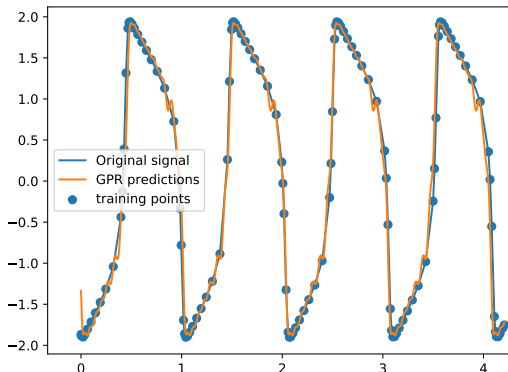
---

## This week's goals

- ✂ Rederive GPR for vector outputs
  - ▶ Turns out this is an open problem
  - ▶ Hard to do generally, but I've found a way to avoid needing this
- ✂ Get GPR to work
  - ▶ Some success
- ✂ Use it for a predictor-corrector
  - ▶ Not got this far yet

---

## Model fitting – it sort of works!



---

## 30 seconds intro to Gaussian processes

- ✂ Gives us a Gaussian distribution at any given time
- ✂ That Gaussian distribution is the PDF of our function value at that time
- ✂ Works by maintaining a probability distribution over candidate functions
- ✂ Constructed by using Bayes' rule to condition on the evidence and form a posterior function distribution
- ✂ Bayes' rule needs good priors!

---

## GPR kernels

- ✿ They specify our prior distributions over functions
- ✿ Good kernels = good priors = good results
- ✿ (Kernels are interesting – they implicitly encode an infinite dimensional feature space)

---

## Periodic kernels

Data are periodic, so it makes sense to have a kernel that's periodic

- ✂ If we choose a periodic kernel then we're favouring periodic functions in our prior function distribution
- ✂ Periodic kernels give better fits on periodic data!
- ✂ But, if the period isn't specified correctly, they'll give big errors and be harder to optimise. . .

---

## Aperiodic kernels

To avoid period-errors, use an aperiodic kernel and overlay each period's data on top of each other

- ✖ Aperiodic kernels don't encode our prior beliefs about periodicity, so they're not going to give as good a fit to the data
- ✖ But, they have no period component, so they aren't sensitive to errors in the period
- ✖ This means that in practice they can actually still give reasonable fits to the data

---

## Hyperparameters

Both periodic and aperiodic kernels rely on hyperparameters; often. . .

- ✿  $l$ : how similar nearby datapoints are
- ✿  $\sigma_f^2$ : function amplitude
- ✿  $\sigma_n^2$ : noise in the function observations


There's an interplay between kernel choice and hyperparameter selection:

- ✿ well-chosen kernels are easier to fit hyperparameters to; will still give good results with bad hyperparameters
- ✿ bad kernels give bad results unless the hyperparameters are perfect, which is hard!



---


## Characteristic lengths



Bigger version on the next slide


- ✦  $l$  is the most interesting hyperparameter
- ✦ Measures how similar near-by datapoints are to each other
- ✦ Since neurons are a multiple-timescale system, this isn't trivial
- ✦ RED: points sampled close in time map to very different values, and are therefore dissimilar; small  $l$
- ✦ GREEN: points sampled close in time map to similar values; big  $l$

# Characteristic lengths



---

## The effects of $l$



✖  $l$  varies across the signal

✖ modelling with constant  $l$  gives bad results

---

## Solution

There's kernels for modelling variable  $l$ , but...

- ✖  $l$  becomes a function in space
- ✖ No longer a single hyperparameter to fit, but an entire hyperfunction
- ✖ Hyperparameter space goes from 3-dimensional to infinite!
- ✖ One approach models the  $l$  function as a Gaussian process, and demonstrates an efficient / computationally tractable way of fitting it
- ✖ The paper is hard

---

## Generalised spectral mixture kernels

- ✿ Use GPR to generate a kernel for the specific input data
- ✿ Provides a tractable way of fitting this kernel
- ✿ Once fitted for one periodic orbit, it will still work well for the rest
- ✿ Automatically deals with periodicity, non-stationarity, so we resolve the periodic kernel dilemma!

Remes, Sami, Markus Heinonen, and Samuel Kaski. "Non-stationary spectral kernels." Advances in Neural Information Processing Systems. 2017.

---

## Next steps

- ✂ Work through the paper to understanding
  - ▶ Might take a while!
- ✂ Implement a GSMKernel
  - ▶ This should finish off the the GPR part
  - ▶ If GPR turns out to be a no-go, the rest of the predictor/corrector scheme will still work with another interpolating model, eg. periodic splines
- ✂ Code up a predictor
  - ▶ Should be trivial once GPR is sorted
- ✂ Code up a corrector
  - ▶ Should be interesting but very doable