

GPR Kernels

Mark Blyth

This week's goals

- ✂ Rederive GPR for vector outputs
 - ▶ Turns out this is an open problem
 - ▶ Hard to do generally, but I've found a way to avoid needing this
- ✂ Get GPR to work
 - ▶ Some success
- ✂ Use it for a predictor-corrector
 - ▶ Not got this far yet

Last time...

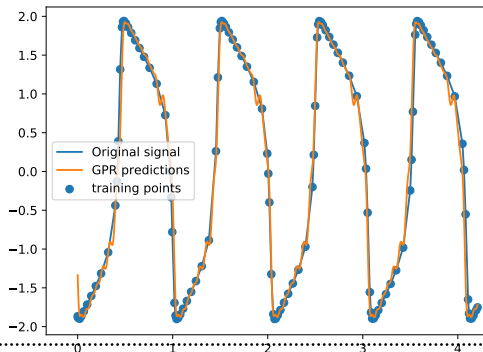
Period windowing (1)

- ✳ We have a periodic signal $f(t')$, taken from our observed system output (here, neuron spikes)
- ✳ We wish to split it into windows $f_1(t), f_2(t), \dots, t \in [0, 1]$, such that $f_i(1) = f_{i+1}(0)$ (periodicity)
- ✳ Then $f_i(t)$ is a function representing the i 'th period of the signal
 - ▶ Eg. if $f(t') = \sin(\frac{t'}{2\pi})$ with $t \in [0, \infty)$, then $f_1(t) = \sin(\frac{t}{2\pi})$, $f_2(t) = \sin(\frac{t}{2\pi} + 2\pi)$, $f_3(t) = \sin(\frac{t}{2\pi} + 4\pi)$, \dots , with $t \in [0, 1]$
 - ▶ By periodicity we have $f_i(t) = f_j(t)$, and $f_i(1) = f_{i+1}(0)$
- ✳ Fitting a model $t \rightarrow f_i(t)$ to these function observations gives us the periodic orbit model $f^*(t)$ at the current parameter value
- ✳ It's hard to split data up into these periods!

Splitting data into periods turns out to be easy:

- ✳ Rescale time as $t = t'/T \bmod T$
- ✳ Periodic kernels take care of this automatically
 - ▶ Choosing an appropriate kernel can make life much easier

Model fitting – it sort of works!



- ✦ Trying to model a Fitzhugh-Nagumo output with a Gaussian process
- ✦ Some wiggleness between datapoints (more on that later!) but it generally works

30 seconds intro to Gaussian processes

- ✂ Gives us a Gaussian distribution at any given time
- ✂ That Gaussian distribution is the PDF of our function value at that time
- ✂ Works by maintaining a probability distribution over candidate functions
- ✂ Constructed by using Bayes' rule to condition on the evidence and form a posterior function distribution
- ✂ Bayes' rule needs good priors!

GPR kernels

- ✿ They specify our prior distributions over functions
- ✿ Good kernels = good priors = good results
- ✿ (Kernels are interesting – they implicitly encode an infinite dimensional feature space)

Periodic kernels

Data are periodic, so it makes sense to have a kernel that's periodic

- ✦ If we choose a periodic kernel then we're favouring periodic functions in our prior function distribution
- ✦ Periodic kernels give better fits on periodic data!
- ✦ But, if the period isn't specified correctly, they'll give big errors and be harder to optimise. . .

Aperiodic kernels

To avoid period-errors, use an aperiodic kernel and overlay each period's data on top of each other

- ✖ Aperiodic kernels don't encode our prior beliefs about periodicity, so they're not going to give as good a fit to the data
- ✖ But, they have no period component, so they aren't sensitive to errors in the period
- ✖ This means that in practice they can actually still give reasonable fits to the data

Hyperparameters


Both periodic and aperiodic kernels rely on hyperparameters; often. . .

- ✿ l : how similar nearby datapoints are
- ✿ σ_f^2 : function amplitude
- ✿ σ_n^2 : noise in the function observations

There's an interplay between kernel choice and hyperparameter selection:

- ✿ well-chosen kernels are easier to fit hyperparameters to; will still give good results with bad hyperparameters
- ✿ bad kernels give bad results unless the hyperparameters are perfect, which is hard!


Characteristic lengths




Bigger version on the next slide

- ✦ l is the most interesting hyperparameter
- ✦ Measures how similar near-by datapoints are to each other
- ✦ Since neurons are a multiple-timescale system, this isn't trivial
- ✦ RED: points sampled close in time map to very different values, and are therefore dissimilar; small l
- ✦ GREEN: points sampled close in time map to similar values; big l

Characteristic lengths



The effects of l



✖ l varies across the signal

✖ modelling with constant l gives bad results

Solution

There's kernels for modelling variable l , but...

- ✿ l itself becomes a function of input variables
- ✿ No longer a single hyperparameter to fit, but an entire hyperfunction
- ✿ Hyperparameter space goes from 3-dimensional to infinite!
- ✿ One approach models the l function as a Gaussian process, and demonstrates an efficient / computationally tractable way of fitting it

Generalised spectral mixture kernels

- ✿ Use GPR to generate a kernel for the specific input data
- ✿ Provides a tractable way of fitting this kernel
- ✿ Once fitted for one periodic orbit, it will still work well for the rest
- ✿ Automatically deals with periodicity, non-stationarity, so we resolve the periodic kernel dilemma!
- ✿ The paper is hard

Remes, Sami, Markus Heinonen, and Samuel Kaski. "Non-stationary spectral kernels." Advances in Neural Information Processing Systems. 2017.

Next steps

- ✂ Work through the paper to understanding
 - ▶ Might take a while!
- ✂ Implement a GSMKernel
 - ▶ This should finish off the the GPR part
 - ▶ If GPR turns out to be a no-go, the rest of the predictor/corrector scheme will still work with another interpolating model, eg. periodic splines
- ✂ Code up a predictor
 - ▶ Should be trivial once GPR is sorted
- ✂ Code up a corrector
 - ▶ Should be interesting but very doable