

# GPR Kernels

Mark Blyth

---

## This week's goals

- ✦ Rederive GPR for vector outputs
  - ▶ Turns out this is an open problem
  - ▶ Hard to do generally, but I've found a way to avoid needing this
- ✦ Get GPR to work
  - ▶ Some success
- ✦ Use it for a predictor-corrector
  - ▶ Not got this far yet

---

## Last time...

---


### Period windowing (1)

- ✦ We have a periodic signal  $f(t')$ , taken from our observed system output (here, neuron spikes)
- ✦ We wish to split it into windows  $f_1(t), f_2(t), \dots, t \in [0, 1]$ , such that  $f_i(1) = f_{i+1}(0)$  (periodicity)
- ✦ Then  $f_i(t)$  is a function representing the  $i$ 'th period of the signal
  - ▶ Eg. if  $f(t') = \sin(\frac{t'}{2\pi})$  with  $t \in [0, \infty)$ , then  $f_1(t) = \sin(\frac{t}{2\pi})$ ,  $f_2(t) = \sin(\frac{t}{2\pi} + 2\pi)$ ,  $f_3(t) = \sin(\frac{t}{2\pi} + 4\pi)$ ,  $\dots$ , with  $t \in [0, 1]$
  - ▶ By periodicity we have  $f_i(t) = f_j(t)$ , and  $f_i(1) = f_{i+1}(0)$
- ✦ Fitting a model  $t \rightarrow f_i(t)$  to these function observations gives us the periodic orbit model  $f^*(t)$  at the current parameter value
- ✦ It's hard to split data up into these periods!

Splitting data into periods turns out to be easy:

- ✦ Rescale time as  $t = t'/T \bmod T$
- ✦ Periodic kernels take care of this automatically
  - ▶ Choosing an appropriate kernel can make life much easier

## Model fitting – it sort of works!



- ✦ Trying to model a Fitzhugh-Nagumo output with a Gaussian process
- ✦ Some wiggleness between datapoints (more on that later!) but it generally works

---

## 30 seconds intro to Gaussian processes

- ✿ Here, takes time as an input, and outputs a Gaussian distribution at that time
- ✿ The Gaussian distribution is the probability distribution of our function value, at that time
- ✿ This works by maintaining a probability distribution over candidate functions
- ✿ Bayes' rule is to condition on the evidence, and form a posterior function distribution
- ✿ *Bayes' rule needs good priors!*

---

## GPR kernels

- ✿ They specify our prior distributions over functions
- ✿ Good kernels = good priors = good results
- ✿ (Kernels are interesting – they implicitly encode an infinite dimensional feature space)

---

## Periodic kernels

Data are periodic, so it makes sense to have a kernel that's periodic

- ✦ If we choose a periodic kernel then we're favouring periodic functions in our prior function distribution
- ✦ Periodic kernels give better fits on periodic data!
- ✦ But, if the period isn't specified correctly, they'll give big errors and be harder to optimise. . .

---

## Aperiodic kernels

To avoid period-errors, use an aperiodic kernel and overlay each period's data on top of each other

- ✖ Aperiodic kernels don't encode our prior beliefs about periodicity, so they're not going to give as good a fit to the data
- ✖ But, they have no period component, so they aren't sensitive to errors in the period
- ✖ This means that in practice they can actually still give reasonable fits to the data



---

## Hyperparameters

Both periodic and aperiodic kernels rely on hyperparameters; often. . .


- ✿  $l$ : how similar nearby datapoints are
- ✿  $\sigma_f^2$ : function amplitude
- ✿  $\sigma_n^2$ : noise in the function observations

There's an interplay between kernel choice and hyperparameter selection:

- ✿ well-chosen kernels are easier to fit hyperparameters to; will still give good results with bad hyperparameters
- ✿ bad kernels give bad results unless the hyperparameters are perfect, which is hard!

---


## Characteristic lengths




Bigger version on the next slide

- ✦  $l$  is the most interesting hyperparameter
- ✦ Measures how similar near-by datapoints are to each other
- ✦ Since neurons are a multiple-timescale system, this isn't trivial
- ✦ RED: points sampled close in time map to very different values, and are therefore dissimilar; small  $l$
- ✦ GREEN: points sampled close in time map to similar values; big  $l$

# Characteristic lengths



## The effects of $l$



✖  $l$  varies across the signal

✖ modelling with constant  $l$  gives bad results

---

## Solution

There's kernels for modelling variable  $l$ , but...

- ✿  $l$  itself becomes a function of input variables
- ✿ No longer a single hyperparameter to fit, but an entire hyperfunction
- ✿ Hyperparameter space goes from 3-dimensional to infinite!
- ✿ One approach models the  $l$  function as a Gaussian process, and demonstrates an efficient / computationally tractable way of fitting it

---

## Generalised spectral mixture kernels

- ✂ Use GPR to generate a kernel for the specific input data
- ✂ Provides a tractable way of fitting this kernel
- ✂ Once fitted for one periodic orbit, it will still work well for the rest
- ✂ Automatically deals with periodicity, non-stationarity, so we resolve the periodic kernel dilemma!
- ✂ The paper is hard

Remes, Sami, Markus Heinonen, and Samuel Kaski. "Non-stationary spectral kernels." Advances in Neural Information Processing Systems. 2017.

---

## Next steps

- ✂ Work through the paper to understanding
  - ▶ Might take a while!
- ✂ Implement a GSMKernel
  - ▶ This should finish off the the GPR part
  - ▶ If GPR turns out to be a no-go, the rest of the predictor/corrector scheme will still work with another interpolating model, eg. periodic splines
- ✂ Code up a predictor
  - ▶ Should be trivial once GPR is sorted
- ✂ Code up a corrector
  - ▶ Should be interesting but very doable