

# PRESENTATION NOTES, DON'T DISPLAY THESE!

Mark Blyth

# Last meeting

Last time:

- ✿ The challenges of working with spiking signals
  - ▶ Lots of high-frequency energy
  - ▶ Typically noisy, both from stochastics and from measurement errors
  - ▶ Lots of high-frequency components mean a LP filter will remove the signal, as well as noise, so we can't naively clean up the signal
  - ▶ Also means using a truncated Fourier discretisation will be infeasible, since it'll have far too many components to effectively discretise
- ✿ Surrogate methods to overcome these:
  - ▶ Use a regression model in place of the real data
  - ▶ Perform desired analysis on this instead
  - ▶ 'Desired analysis' will be explained more later, in a CBC context
  - ▶ A well-chosen surrogate will filter out all the noise, without losing any signal

# Surrogates point 1

✿ Given  $y_i = f(x_i) + \varepsilon$ , estimate  $f(x)$

- ▶ We assume there's a 'true' underlying signal  $f(x)$
- ▶ This true signal is what the neuron is actually doing, eg. what the membrane potential actually is at the patch clamp location
- ▶ We don't have access to  $f(x)$ ; instead, we get a time-series  $y_i$ , of noise-corrupted samples
- ▶ These noise-corrupted samples contain both the actual signal at the given sample time, plus some nuisance variable  $\varepsilon$  from errors in measurement
- ▶ We wish to recover  $f(x)$  from these samples, as that's the noise-free, true signal that we're interested in
  - ▶ Simply LP filtering would only remove the HF components of  $f(x)$  and  $\varepsilon$ ; instead, we wish to separate the two out into noise and signal
- ▶ We can use statistical methods to infer  $f(x)$  and  $\varepsilon$
- ▶ This gives us a clean, noise-free surrogate to perform all the analysis on
- ▶ Surrogate: we use  $f(x)$  in place of the real data

# Surrogates point 2

✂ Splines and Gaussian processes are good methods for estimation

- ▶ GPR: mathematically elegant, rigorous method
  - ▶ Assume normally distributed  $\varepsilon$
  - ▶ Point estimate:  $y \sim N(f(x), \varepsilon)$
  - ▶ Whole function estimate:  $f(x) \sim GP(\mu, \Sigma)$
  - ▶ Whole function estimate is actually a Gaussian distribution over functions
  - ▶ The whole function estimate is just a generalisation of the point estimate
  - ▶ For a sensible prior, we can then use Bayes to estimate the posterior distribution on  $f(x)$
  - ▶ Statistically optimal estimator
  - ▶ Downside: for finite data, results are only as good as the priors we use; coming up with good priors is hard
- ▶ Splines: simple, effective, less elegant
  - ▶ Split  $f(x)$  into intervals, and assume  $f(x)$  is locally polynomial on any given interval
  - ▶ Enforce  $C^2$  smoothness over polynomial sections
  - ▶ Polynomials then join up at the edges of each interval; these joinings are called knot points
  - ▶ Remaining free parameters are chosen to maximise goodness-of-fit
  - ▶ No need to define priors, so it's easier to use on data where choice of priors becomes difficult
  - ▶ Knot points are difficult to choose; use Bayesian inference to form a posterior distribution over knots

# Developments since last time

## IMAGE.

### ✂ Surrogates tested and working

- ▶ GPR works on both real and synthetic data, in cases where the data are sufficiently stationary
- ▶ Free knot splines works always, so use it in cases where the data aren't sufficiently stationary

### ✂ Image taken from recent abstract

- ▶ Bayesian free-knot splines
- ▶ Works well – we can extract the underlying signal near perfectly, even given very noisy observations
- ▶ Three changepoints per period: at the start, top, and end of a spike; here, the signal rapidly changes from slow to fast behaviour
- ▶ These changepoints are the hardest bits to model, and therefore the surrogates are least accurate here (the time between spikes shows a slow, gentle change that's easy to model accurately)
- ▶ Zooming in on one of the changepoints, we see that the surrogate recreates the latent signal nearly exactly; even at the most difficult-to-fit part of the signal, we still get excellent results

# Developments since last time

## ✶ Novel discretisations

- ▶ Surrogates only give us a noise filter
- ▶ For some CBC implementations, this is sufficient
- ▶ In CBC cases where we have to do Newton iterations, this isn't useful; we instead need a low-dimensional discretisation
- ▶ We can apply the surrogates ideas to creating discretisations

## ✶ In-silico CBC

- ▶ Best way to demonstrate that these methods work, are valuable

# Discretisations

- ✂ Discretisation takes a function, projects it onto a set of basis functions
- ✂ Coefficients and basis functions are sufficient to represent the signal
- ✂ Lots of possible choices for basis functions
  - ▶  $C^{\infty}$  signals can be represented exactly with monomial basis functions (taylor expansion)
  - ▶ Periodic signals can be represented exactly with trig basis functions (Fourier series)
  - ▶ These are bad choices for neuron CBC – require lots of coefficients to describe the spiking signals
- ✂ We've already met splines; turns out we can define a set of basis functions for splines
  - ▶ Can therefore express any spline curve in the above form
  - ▶ This means we can discretise with splines too!
  - ▶ Splines are a good choice: they provide a nice simple, intuitive model, and don't require many basis functions to get a good approximation

# Splines discretisation

- ✂ Fit a set of basis functions to initial signal  $f_0(x)$ 
  - ▶ Choose a set of knots  $x_i$ , such that the splines basis  $b_i(x)$  that we construct from knots  $x_i$  is able to fit the initial signal  $f_0(x)$  as accurately as possible, in the least squares sense
  - ▶ This is actually hard to do – open research problem
  - ▶ Elegant approach: find a Bayesian posterior over  $\xi|data$ . Downside: is slow and complicated; need to do MCMC to approximate intractable integral
  - ▶ Simple approach 1: put a knot at every datapoint then penalise functional of second derivative, to enforce smoothness. Downside: we end up with huge numbers of knots.
  - ▶ Simple approach 2: keep adding knots until we reach satisfactory results; downside: lower quality fit, no guarantee of low-dimensionality
- ✂ My approach: choose the number of knots; numerically optimise knot positions; start from random initial knots; avoid local minima by repeating this lots
  - ▶ Downside: need to repeat lots to find global minimum
  - ▶ Need to choose the number of knots a priori; algo doesn't work it out for us
  - ▶ Upside: quick and easy approach to finding a good set of knots; easiest way to get low-dimensional knot set



# Splines demo

- ✂ Splines discretisation works well
- ✂ This example uses just 8 knot points
  - ▶ Higher than an 8d discretisation, as we need to add exterior knots so that the basis splines have support across the range of the data
- ✂ Reconstructs the latent signal near-perfectly

# Splines vs Fourier

Also shown: Fourier

- ✖ Visually, splines fits better than Fourier
- ✖ Fourier is harder to fit
- ✖ Too few harmonics and the series can't fit the data
- ✖ Too many harmonics and the series starts fitting the noise as well as the data
- ✖ Not really any sweet spot; no point where the series fits the signal, but averages out the noise
- ✖ This is the usage case for surrogates – when we have noisy data, but still want to use Fourier with it!

# Goodness-of-fit

This shows the goodness-of-fit of a splines model with given number of knots, and Fourier series with given number of harmonics

- ✂ No noise, Fitzhugh Nagumo
- ✂ Splines error decays more rapidly than Fourier error
- ✂ Effects become even more dramatic for more neuron-like signals
- ✂ Note though this is the goodness-of-fit of a splines, fourier model on a single signal; doesn't determine how well the splines model generalises to discretising unseen signals, ie. only shows how well the spline model fits a signal to which its basis functions were fitted; using the same basis functions on a signal from a different parameter value might get different goodness-of-fit. Fourier won't have this issue since it uses trig basis all the time

# Method usage cases

## ✶ Harmonically forced:

- ▶ When we have a harmonically forced system, we can have a harmonically oscillating control action, and treat the control action as the forcing term
- ▶ In this setup, we can efficiently iterate on the Fourier harmonics, to drive the higher-order harmonics of the control action to zero
- ▶ This necessitates a Fourier projection. No need for a novel discretisation, but we could possibly improve the Fourier discretisation by using a surrogate to first filter off the noise

## ✶ Non-harmonically forced:

- ▶ If system is unforced, we apply parameter and control action separately, and need the control action to be zero
- ▶ We can use Newton iterations to solve for the noninvasive control action
- ▶ Since we're doing Newton iterations, we need to work with a low dimensional system, otherwise it'll be impractically slow
- ▶ To have a low-dimensional system, we use a novel discretisation, eg. splines

# *In-silico* CBC

Current work: implementing an in-silico CBC simulation

- ✶ Best way to test if the surrogates, discretisations work with CBC is to try using them with CBC!

# CBC method POINTS 1, 2

- ✂ Use PD control
  - ▶ Easy, model-free control method
  - ▶ Gets good results with a method we could easily use in experiments too
  - ▶ Fit control parameters with brute force
  - ▶ Easy to simulate, minimal effort in controller design
- ✂ As per standard numerical continuation, do a change in variables so that time is in  $[0,1]$ , and treat period as an extra continuation variable
  - ▶ Not necessary with Fourier discretisation
  - ▶ Splines knots are like finite-differences or collocation mesh points
  - ▶ Time rescaling is necessary with mesh-based methods, as changing the period would effectively move the mesh points relative to the signal

# CBC method POINTS 3, 4

## ✂ Non-adaptive mesh

- ▶ Fit knots at the start, keep them in the same position throughout
- ▶ Adaptive mesh would mean re-fitting the knots after a prediction-correction step
- ▶ In terms of code, this is minimal extra effort, but would add a slight fitting overhead
- ▶ Only using non-adaptive mesh because I'm interested to see how well it works

## ✂ Use Newton iterations to solve for discretised control target = discretised system output

- ▶ Nothing fancy, just simple, slow Newton with finite differences; able to do this in-silico, but would need more rigorous treatment for experiments
- ▶ Sensible to start off with easy root finding, and develop something fancy (Broyden) later
- ▶ Splines method adds exterior knots, and some of the coefficients are always zero, so they can be removed from the discretisation to speed up the finite-differences Jacobian step; I'm currently being lazy and not doing this

# Discretisors

- ✂ Instantiate desired discretisor type with its relevant parameter
  - ▶ Fourier:  $n_{\text{harmonics}}$
  - ▶ Splines: knot locations
  - ▶ Regardless of discretisor type, can then call discretisor discretise, discretisor undiscretise
- ✂ Simple, standard interface to discretisation routines
  - ▶ Able to swap between Fourier, Splines with zero effort
  - ▶ Allows direct comparison between discretisation methods
  - ▶ Could easily implement any other discretisation (eg. wavelets) using the same interface
- ✂ Lightly tested:
  - ▶ The code runs and produces sensible outputs
  - ▶ Haven't tested its ability to generalise to new signals
  - ▶ I.e. don't know how well basis funcs fitted to  $f_0(x)$  will work for discretising  $f_1(x)$



# Controllers

- ✶ Can design controllers with a standard interface too
  - ▶ Set the controller type, control target, gains, and the control matrices
  - ▶ The controller object handles the rest
- ✶ Similarly, can design systems with a standard interface
  - ▶ Specify a function that gives the ODE RHS; a list of ODE parameters; the controller
  - ▶ Can then run the controlled model for any choice of time range, ICs, pars
  - ▶ Subsequent runs optionally start with ICs given by final state of last run, much like a real system

The point of all these standardised interfaces is that it becomes really easy to swap everything out; eg. apply to different models, different control strategies, different discretisers

Can then run a CBC experiment in less than 10 lines of code; easy to apply, reapply, experiment with

# Control

## IMAGE

FH system with sine target; looks very reasonable

- ✂ Lightly tested: code runs, results look very reasonable
  - ▶ Seems like a sensible output
- ✂ Can easily write out the RHS of a PD-controlled FH system w/ sine target; can compare this explicit RHS to the code-generated system, to make sure the code isn't doing anything funny
  - ▶ Haven't done this yet

# Continuation

- ✚ Runs a psuedo-arclength secant-predictor Newton-corrector CBC
- ✚ Code is written now
- ✚ Requires a 'system': this is just a controlled model [from previous code], with its arguments binded
- ✚ 'system' interface will be easy to write, but haven't got round to this yet

# Simulation summary

- ✂ Results handling is just something to take the set of natural periodic orbits, apply some measure (eg. amplitude), then plot them on a bifurcation diagram.
- ✂ Should be working and tested within a week

# Open questions

## ✂ Will splines discretisation work?

- ▶ If splines can only model the signal to which the knots were fitted, they won't work for CBC
- ▶ My guess is they will work

## ✂ Stationary or adaptive mesh?

- ▶ If splines basis aren't good at generalising, can re-fit knots at each step, much like an adaptive mesh, which would hopefully fix problems

## ✂ Efficient solving methods

- ▶ Remove zero-coefficients from discretisation
- ▶ Broyden Jacobian update?
- ▶ Newton-Picard iterations? Ludovic's suggestion of Newton-iterating on unstable coefficients, fixed-point iterating on stable coefficients; reduces the size of the Jacobian / finite differences step

## ✂ Can we interface the code with Simulink?

- ▶ Ludovic has a simulink model that would be fun to play with; haven't looked at it yet since I've been testing the CBC codes; would be interesting to try to call the finished code from MATLAB, in which case we might be able to interface the two

# Next steps

- ✂ Continuation tutorial paper
  - ▶ Haven't touched it recently
  - ▶ Making slower progress since I'm trying to get the stuff for this done before the paper deadline
- ✂ NODYCON abstract for this: submitted, accepted
- ✂ Conference paper for this: will start on that once the CBC simulation is sorted