

Bayesian free-knot splines

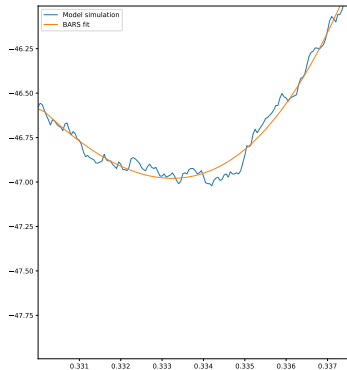
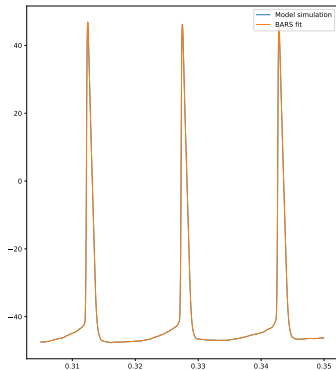
Mark Blyth

Week's goals

- ✂ Make changes to continuations paper
 - ▶ Looked at feedback, haven't started making changes yet
- ✂ Fix MSPE downsampling errors
 - ▶ Haven't fixed this yet
- ✂ Implement and test free-knot splines
 - ▶ Learn how it works
 - ▶ Code it up
 - ▶ Use it to validate splines method

Bayesian free-knot splines

Identified as being a good method for modelling neuron data



How it works

1. Assume a spline model fits the data
2. Find a distribution over spline models, given the data
3. Condition on this distribution with new data, to get posterior estimates

Step 0: problem setup

✿ Take some data $Y_i = g(x_i) + \varepsilon$

- ▶ $\varepsilon \sim \mathcal{N}\{0, \sigma^2\}$
- ▶ σ unknown
- ▶ g unknown
- ▶ Y_i random variables
- ▶ Then, $Y_i | x_i, \sigma \sim \mathcal{N}\{g(x_i), \sigma^2\}$

✿ Goal: estimate g from noisy samples (x_i, Y_i)

This is a very standard problem formulation so far...

Step 1: spline model

- ✂ Model latent function g as being some piecewise-polynomial function f
- ✂ Tie polynomials together at knot-points ξ_i

$$f(x) = \begin{cases} f_1(x), & x \in [a, \xi_0) \\ f_2(x), & x \in [\xi_0, \xi_1) \\ \dots \\ f_{k+2}, & x \in [\xi_k, b] \end{cases} \quad (1)$$

- ✂ $f_i(x)$ is an $\mathcal{O}(3)$ polynomial passing through $(\xi_{i-1}, g(\xi_{i-1}))$, $(\xi_i, g(\xi_i))$
 - ... or allow the polynomials to pass *near* the knot-points, for smoothing splines

Spline bases

- ✦ Calculating the coefficients for each f_i is inconvenient

Spline bases

- ✎ Calculating the coefficients for each f_i is inconvenient
- ✎ Nicer approach:

Spline bases

- ✂ Calculating the coefficients for each f_i is inconvenient
- ✂ Nicer approach:
 - ▶ Let $f(x) = \sum_i \beta_i b_i(x)$

Spline bases

- ✂ Calculating the coefficients for each f_i is inconvenient
- ✂ Nicer approach:
 - ▶ Let $f(x) = \sum_i \beta_i b_i(x)$
 - ▶ Functions b_i form a basis for some spline $f(x)$ with k knots at $\{\xi_0, \dots, \xi_k\}$

Spline bases

- ✂ Calculating the coefficients for each f_i is inconvenient
- ✂ Nicer approach:
 - ▶ Let $f(x) = \sum_i \beta_i b_i(x)$
 - ▶ Functions b_i form a basis for some spline $f(x)$ with k knots at $\{\xi_0, \dots, \xi_k\}$
- ✂ b_i are found by

Spline bases

- ✂ Calculating the coefficients for each f_i is inconvenient
- ✂ Nicer approach:
 - ▶ Let $f(x) = \sum_i \beta_i b_i(x)$
 - ▶ Functions b_i form a basis for some spline $f(x)$ with k knots at $\{\xi_0, \dots, \xi_k\}$
- ✂ b_i are found by
 - ▶ Specifying knot locations

Spline bases

- ✂ Calculating the coefficients for each f_i is inconvenient
- ✂ Nicer approach:
 - ▶ Let $f(x) = \sum_i \beta_i b_i(x)$
 - ▶ Functions b_i form a basis for some spline $f(x)$ with k knots at $\{\xi_0, \dots, \xi_k\}$
- ✂ b_i are found by
 - ▶ Specifying knot locations
 - ▶ Requiring \mathcal{C}^1 smoothness

Spline bases

- ✂ Calculating the coefficients for each f_i is inconvenient
- ✂ Nicer approach:
 - ▶ Let $f(x) = \sum_i \beta_i b_i(x)$
 - ▶ Functions b_i form a basis for some spline $f(x)$ with k knots at $\{\xi_0, \dots, \xi_k\}$
- ✂ b_i are found by
 - ▶ Specifying knot locations
 - ▶ Requiring \mathcal{C}^1 smoothness
 - ▶ Assuming linearity outside of data range

Spline bases

- ✂ Calculating the coefficients for each f_i is inconvenient
- ✂ Nicer approach:
 - ▶ Let $f(x) = \sum_i \beta_i b_i(x)$
 - ▶ Functions b_i form a basis for some spline $f(x)$ with k knots at $\{\xi_0, \dots, \xi_k\}$
- ✂ b_i are found by
 - ▶ Specifying knot locations
 - ▶ Requiring \mathcal{C}^1 smoothness
 - ▶ Assuming linearity outside of data range
- ✂ b_i are called basis splines

Spline bases

- ✂ Calculating the coefficients for each f_i is inconvenient
- ✂ Nicer approach:
 - ▶ Let $f(x) = \sum_i \beta_i b_i(x)$
 - ▶ Functions b_i form a basis for some spline $f(x)$ with k knots at $\{\xi_0, \dots, \xi_k\}$
- ✂ b_i are found by
 - ▶ Specifying knot locations
 - ▶ Requiring \mathcal{C}^1 smoothness
 - ▶ Assuming linearity outside of data range
- ✂ b_i are called basis splines
 - ▶ Our model now becomes $Y_i | x_i, \beta, \sigma, \xi \sim \mathcal{N}\{\sum_i \beta_i b_i(x_i), \sigma^2\}$

Easy approach

- ✦ Choose a nice number of knots k

Downside: bad choices for any of these parameters will give bad results:

Easy approach

- ✦ Choose a nice number of knots k
- ✦ Choose a uniformly spaced knot-set ξ

Downside: bad choices for any of these parameters will give bad results:

Easy approach

- ✦ Choose a nice number of knots k
- ✦ Choose a uniformly spaced knot-set ξ
- ✦ Guess σ

Downside: bad choices for any of these parameters will give bad results:

Easy approach

- ✦ Choose a nice number of knots k
- ✦ Choose a uniformly spaced knot-set ξ
- ✦ Guess σ
- ✦ Find a MLE for β

Downside: bad choices for any of these parameters will give bad results:

Easy approach

- ✂ Choose a nice number of knots k
- ✂ Choose a uniformly spaced knot-set ξ
- ✂ Guess σ
- ✂ Find a MLE for β

Downside: bad choices for any of these parameters will give bad results:

- ✂ Too few knots = underparameterised = can't capture shape of data

Easy approach

- ✂ Choose a nice number of knots k
- ✂ Choose a uniformly spaced knot-set ξ
- ✂ Guess σ
- ✂ Find a MLE for β

Downside: bad choices for any of these parameters will give bad results:

- ✂ Too few knots = underparameterised = can't capture shape of data
- ✂ Too many knots = overparameterised = overfit data and capture noise

Easy approach

- ✂ Choose a nice number of knots k
- ✂ Choose a uniformly spaced knot-set ξ
- ✂ Guess σ
- ✂ Find a MLE for β

Downside: bad choices for any of these parameters will give bad results:

- ✂ Too few knots = underparameterised = can't capture shape of data
- ✂ Too many knots = overparameterised = overfit data and capture noise
- ✂ Bad knot positioning = function can't adapt to changing rates

Step 2: find a distribution over models

- ✦ Specify a prior belief $\pi_k(k)$ for the number of knots we have

Joint probability: $p(k, \xi, \beta, \sigma, y) = p(y|\beta, \sigma)\pi_\sigma(\sigma)\pi_\beta(\beta|\sigma, \xi, k)\pi_\xi(\xi|k)\pi_k(k)$

We can evaluate all of this!

Step 2: find a distribution over models

- ✦ Specify a prior belief $\pi_k(k)$ for the number of knots we have
 - ▶ Eg. discrete uniform on $[k_{min}, k_{max}]$

Joint probability: $p(k, \xi, \beta, \sigma, y) = p(y|\beta, \sigma)\pi_\sigma(\sigma)\pi_\beta(\beta|\sigma, \xi, k)\pi_\xi(\xi|k)\pi_k(k)$

We can evaluate all of this!

Step 2: find a distribution over models

- ✦ Specify a prior belief $\pi_k(k)$ for the number of knots we have
 - ▶ Eg. discrete uniform on $[k_{min}, k_{max}]$
- ✦ Specify a prior belief $\pi_\xi(\xi|k)$ on the knot positions ξ , for any given number of knots

Joint probability: $p(k, \xi, \beta, \sigma, y) = p(y|\beta, \sigma)\pi_\sigma(\sigma)\pi_\beta(\beta|\sigma, \xi, k)\pi_\xi(\xi|k)\pi_k(k)$

We can evaluate all of this!

Step 2: find a distribution over models

- ✎ Specify a prior belief $\pi_k(k)$ for the number of knots we have
 - ▶ Eg. discrete uniform on $[k_{min}, k_{max}]$
- ✎ Specify a prior belief $\pi_\xi(\xi|k)$ on the knot positions ξ , for any given number of knots
 - ▶ Eg. uniform on range of data

Joint probability: $p(k, \xi, \beta, \sigma, y) = p(y|\beta, \sigma)\pi_\sigma(\sigma)\pi_\beta(\beta|\sigma, \xi, k)\pi_\xi(\xi|k)\pi_k(k)$

We can evaluate all of this!

Step 2: find a distribution over models

- ✿ Specify a prior belief $\pi_k(k)$ for the number of knots we have
 - ▶ Eg. discrete uniform on $[k_{min}, k_{max}]$
- ✿ Specify a prior belief $\pi_\xi(\xi|k)$ on the knot positions ξ , for any given number of knots
 - ▶ Eg. uniform on range of data
- ✿ Specify a prior belief $\pi_\sigma(\sigma)$ on the noise level

Joint probability: $p(k, \xi, \beta, \sigma, y) = p(y|\beta, \sigma)\pi_\sigma(\sigma)\pi_\beta(\beta|\sigma, \xi, k)\pi_\xi(\xi|k)\pi_k(k)$

We can evaluate all of this!

Step 2: find a distribution over models

- ✿ Specify a prior belief $\pi_k(k)$ for the number of knots we have
 - ▶ Eg. discrete uniform on $[k_{min}, k_{max}]$
- ✿ Specify a prior belief $\pi_\xi(\xi|k)$ on the knot positions ξ , for any given number of knots
 - ▶ Eg. uniform on range of data
- ✿ Specify a prior belief $\pi_\sigma(\sigma)$ on the noise level
- ✿ Specify a prior on β

Joint probability: $p(k, \xi, \beta, \sigma, y) = p(y|\beta, \sigma)\pi_\sigma(\sigma)\pi_\beta(\beta|\sigma, \xi, k)\pi_\xi(\xi|k)\pi_k(k)$

We can evaluate all of this!

Bayesian approach

✿ We want to know where to put the knots

✿ Bayesian approach: find the posterior knot distribution $p(k, \xi|y)$

$$p(k, \xi|y) = \frac{p(k, \xi, y)}{p(y)}, \quad (2)$$

$$p(k, \xi, y) = \int \int p(k, \xi, \beta, \sigma, y) d\beta d\sigma \quad (3)$$

$$= \int \int p(y|\beta, \sigma) \pi_\sigma(\sigma) \pi_\beta(\beta|\sigma, \xi, k) \pi_\xi(\xi|k) \pi_k(k) d\sigma d\beta \quad (4)$$

Bayesian approach

Putting it together, we get

$$p(k, \xi | y) = \frac{\int \int p(y | \beta, \sigma) \pi_{\sigma}(\sigma) \pi_{\beta}(\beta | \sigma, \xi, k) \pi_{\xi}(\xi | k) \pi_k(k) d\sigma d\beta}{p(y)} \quad (5)$$

$$= \frac{\int \int p(y | \beta, \sigma) \pi_{\sigma}(\sigma) \pi_{\beta}(\beta | \sigma, \xi, k) \pi_{\xi}(\xi | k) \pi_k(k) d\sigma d\beta}{\sum_k \int \int \int p(k, \xi, \beta, \sigma, y) d\xi d\beta d\sigma} \quad (6)$$

... which is analytically intractable

MCMC sampling

- ✿ Bayesian inference gives posteriors of form

$$\text{posterior} = \frac{\text{likelihood} \times \text{prior}}{\text{Normalising constant}}$$

- ✿ The normalising constant is regularly analytically intractable
- ✿ Markov-chain Monte carlo methods allow us to sample from the posterior distribution anyway

MCMC

MCMC sets up a Markov chain whose stationary distribution is equal to the posterior distribution:

- ✶ Generate a random state from a proposal distribution

MCMC

MCMC sets up a Markov chain whose stationary distribution is equal to the posterior distribution:

- ✦ Generate a random state from a proposal distribution
- ✦ Accept it with some probability

MCMC

MCMC sets up a Markov chain whose stationary distribution is equal to the posterior distribution:

- ✎ Generate a random state from a proposal distribution
- ✎ Accept it with some probability
- ✎ Reject it with some probability

MCMC

MCMC sets up a Markov chain whose stationary distribution is equal to the posterior distribution:

- ✦ Generate a random state from a proposal distribution
- ✦ Accept it with some probability
- ✦ Reject it with some probability
- ✦ On acceptance, change the current state to the accepted state; else, remain at current state

MCMC

MCMC sets up a Markov chain whose stationary distribution is equal to the posterior distribution:

- ✦ Generate a random state from a proposal distribution
- ✦ Accept it with some probability
- ✦ Reject it with some probability
- ✦ On acceptance, change the current state to the accepted state; else, remain at current state
- ✦ Acceptance and rejection probabilities are chosen such that the distribution of accepted states matches that of the prior

MCMC

MCMC sets up a Markov chain whose stationary distribution is equal to the posterior distribution:

- ✂ Generate a random state from a proposal distribution
- ✂ Accept it with some probability
- ✂ Reject it with some probability
- ✂ On acceptance, change the current state to the accepted state; else, remain at current state
- ✂ Acceptance and rejection probabilities are chosen such that the distribution of accepted states matches that of the prior
- ✂ Doesn't require us to calculate the normalising constant!

Reversible-jump MCMC

- ✂ States are the model configuration (k, ξ)
- ✂ These are of many different dimensions
- ✂ To sample from a posterior with varying dimension, we use reversible-jump MCMC
 - ▶ Jump up and down in dimension, probabilistically
 - ▶ Do so in such a way that the posterior is accurate both within and across dimensions

Model inference

- ✂ Using RJMCMC, we can sample from the posterior $p(k, \xi|y)$, even though the dimensionality of ξ is not fixed
- ✂ We can use samples $k, \xi|y$ to condition on new data (x^*, y^*)
 - ▶ $p(y^*|x^*) = p(y^*|k, \xi, x)p(k, \xi|y)$
- ✂ We predict new points without ever actually setting up a splines model
 - ▶ Find a probability distribution over candidate splines models
 - ▶ Weight each spline model's output according to its probability

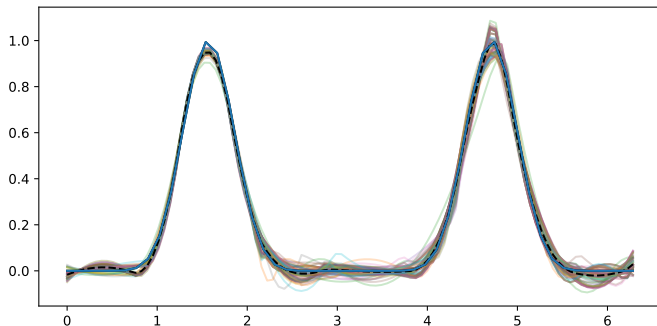
My results

- ✿ Three different MCMC actions can be taken
 - ▶ Add a new knot
 - ▶ Relocate a knot
 - ▶ Delete a knot
- ✿ Each action has a proposal probability (how likely are we to take this action?)
- ✿ Each step has an acceptance probability (how likely are we to accept this action?)
- ✿ The BARS paper does a rather bad job of explaining these!

In my implementation, probabilities are sometimes coming back negative, making it crash

Results

✖ Results can't be trusted!



BARS and GPR

- ✂ BARS maintains a distribution over splines
- ✂ GPR maintains a distribution over arbitrarily many functions
- ✂ Both methods refine the distribution with Bayesian methods

- ✂ BARS probabilistically finds the most informative knot point configuration
 - ▶ Finds set of spline-points that tell us the most about the data
 - ▶ Sparse GPR probabilistically finds the most informative inducing points distribution
- ✂ Tenuous link to optimal experiment design?

Next steps

1. Redraft paper
2. Get BARS to work
 - ▶ Useful as it's the most promising method for a conference abstract
 - ▶ Either get my implementation working, or adapt C code to my needs
3. Fix MSPEs
 - ▶ *Should* be quick and easy
4. (Re)validate all the models I'm playing with
5. Put results into a conference abstract