

# More model fitting

Mark Blyth


---

## Presentation points

- ✦ Models for wind tunnel data
- ✦ A BARS redesign
- ✦ BARS potential pitfalls

---

## An exciting aside




My BARS implementation works!

- ✿ Gives okay results *[some sharp corners?]*
- ✿ Allows arbitrary interpolation
- ✿ Slower than the C implementation

---

## BARS - wind tunnel




2500 datapoints, 41s run-time

---

---


## BARS - wind tunnel



2500 datapoints, 40s run-time


---

## Matern32 (GPR) - wind tunnel



2500 datapoints, 119s run-time

## Matern32 (GPR) - wind tunnel



2500 datapoints, 119s run-time

---


## Timings

- ✂ Unoptimized C code is still faster than optimized python / matlab
  - ▶ BARS is using optimized C implementation, GPRs is using my unoptimized python implementation
  - ▶ GPFlow, GPyTorch would be much faster for GPR
  - ▶ Probabilistic programming might speed up BARS
- ✂ Both methods are  $\mathcal{O}(n^3)$ ; require inverting  $n \times n$  matrix
  - ▶ Both perform poorly on lots of datapoints!



---

## BARS




10,000 datapoints [*crashes with full dataset*], 142s run-time

---

---

## BARS



---

10,000 datapoints [*crashes with full dataset*], max. number of knots [80], more  
burn-in; 590s (9 mins 50s) run-time

---

## BARS issues

- ✖ C implementation doesn't allow validation
  - ▶ Can't interpolate at arbitrary points
  - ▶ Python implementation does! But it's slower
- ✖ C implementation caps the max. number of knots
  - ▶ Python implementation doesn't! But it's slower
  - ▶ Doesn't work for long time series
- ✖  $\mathcal{O}(n^3)$  means data need downsampling to be used
  - ▶ Took about 10 minutes to fit a model to 10% of the data

---

## Experimental data

- ✂ If experiments use small numbers of periods (1, 2, 3), current methods work
  - ▶ Similar to how I tested models on small numbers of neuron spikes
  - ▶ Few periods = less data = fewer knots needed, and quicker fit obtained
- ✂ More datapoints mean we need to get creative...
  - ▶ Sparse variational BARS
  - ▶ Periodic BARS
  - ▶ Semiperiodic methods

Or... speed up BARS by not using BARS

---

## A new BARS approach

IDEA: Bayesian adaptive *evolutionary* splines; genetic algorithm

- ✦ Ideal result: find the MAP knot-config [*best given data*] in fewer steps than MCMC takes to estimate posterior distribution

---

## A new BARS approach

IDEA: Bayesian adaptive *evolutionary* splines; genetic algorithm

- ✎ Ideal result: find the MAP knot-config [*best given data*] in fewer steps than MCMC takes to estimate posterior distribution
  - ▶ BARS uses MCMC to find a posterior knot distribution  $p(k, \xi|y)$

---

## A new BARS approach

IDEA: Bayesian adaptive *evolutionary* splines; genetic algorithm

- ✦ Ideal result: find the MAP knot-config [*best given data*] in fewer steps than MCMC takes to estimate posterior distribution
  - ▶ BARS uses MCMC to find a posterior knot distribution  $p(k, \xi|y)$
  - ▶ MCMC uses 10,000+ knot-shuffling steps to estimate this posterior distribution

---

## A new BARS approach

IDEA: Bayesian adaptive *evolutionary* splines; genetic algorithm

- ✂ Ideal result: find the MAP knot-config [*best given data*] in fewer steps than MCMC takes to estimate posterior distribution
  - ▶ BARS uses MCMC to find a posterior knot distribution  $p(k, \xi|y)$
  - ▶ MCMC uses 10,000+ knot-shuffling steps to estimate this posterior distribution
- ✂ Instead of estimating posterior distribution, why not find posterior mode?



---

## A new BARS approach

IDEA: Bayesian adaptive *evolutionary* splines; genetic algorithm

- ✦ Ideal result: find the MAP knot-config [*best given data*] in fewer steps than MCMC takes to estimate posterior distribution
  - ▶ BARS uses MCMC to find a posterior knot distribution  $p(k, \xi|y)$
  - ▶ MCMC uses 10,000+ knot-shuffling steps to estimate this posterior distribution
- ✦ Instead of estimating posterior distribution, why not find posterior mode?
  - ▶ Use the knot-shuffling steps to evolve an optimal knot-set

---

## A new BARS approach

IDEA: Bayesian adaptive *evolutionary* splines; genetic algorithm

- ✂ Ideal result: find the MAP knot-config [*best given data*] in fewer steps than MCMC takes to estimate posterior distribution
  - ▶ BARS uses MCMC to find a posterior knot distribution  $p(k, \xi|y)$
  - ▶ MCMC uses 10,000+ knot-shuffling steps to estimate this posterior distribution
- ✂ Instead of estimating posterior distribution, why not find posterior mode?
  - ▶ Use the knot-shuffling steps to evolve an optimal knot-set
  - ▶ Find  $\operatorname{argmax}_{k, \xi} p(y|k, \xi)$  [MLE] *or*  $\operatorname{argmax}_{k, \xi} p(k, \xi|y)$  [MAP]

---

## A new BARS approach

IDEA: Bayesian adaptive *evolutionary* splines; genetic algorithm

- ✂ Ideal result: find the MAP knot-config [*best given data*] in fewer steps than MCMC takes to estimate posterior distribution
  - ▶ BARS uses MCMC to find a posterior knot distribution  $p(k, \xi|y)$
  - ▶ MCMC uses 10,000+ knot-shuffling steps to estimate this posterior distribution
- ✂ Instead of estimating posterior distribution, why not find posterior mode?
  - ▶ Use the knot-shuffling steps to evolve an optimal knot-set
  - ▶ Find  $\operatorname{argmax}_{k, \xi} p(y|k, \xi)$  [MLE] *or*  $\operatorname{argmax}_{k, \xi} p(k, \xi|y)$  [MAP]
  - ▶ Each MCMC knot-shuffle becomes a mutation step

---

## A new BARS approach

IDEA: Bayesian adaptive *evolutionary* splines; genetic algorithm

- ✂ Ideal result: find the MAP knot-config [*best given data*] in fewer steps than MCMC takes to estimate posterior distribution
  - ▶ BARS uses MCMC to find a posterior knot distribution  $p(k, \xi|y)$
  - ▶ MCMC uses 10,000+ knot-shuffling steps to estimate this posterior distribution
- ✂ Instead of estimating posterior distribution, why not find posterior mode?
  - ▶ Use the knot-shuffling steps to evolve an optimal knot-set
  - ▶ Find  $\operatorname{argmax}_{k, \xi} p(y|k, \xi)$  [MLE] *or*  $\operatorname{argmax}_{k, \xi} p(k, \xi|y)$  [MAP]
  - ▶ Each MCMC knot-shuffle becomes a mutation step
  - ▶ Each MCMC acceptance/rejection becomes an evolutionary scoring step

---

## Possible evolutionary splines implementation

1. Start with a set of initial (randomly drawn) pool of knot configurations

---

## Possible evolutionary splines implementation

1. Start with a set of initial (randomly drawn) pool of knot configurations
2. For each perturbation step...

---

## Possible evolutionary splines implementation

1. Start with a set of initial (randomly drawn) pool of knot configurations
2. For each perturbation step. . .
  - 2.1 Randomly move, add, or delete a knot from each set

---

## Possible evolutionary splines implementation

1. Start with a set of initial (randomly drawn) pool of knot configurations
2. For each perturbation step. . .
  - 2.1 Randomly move, add, or delete a knot from each set
  - 2.2 If the perturbed knot has a better likelihood  $p(y|k, \xi)$  than the original, replace the original with it



---

## Possible evolutionary splines implementation

1. Start with a set of initial (randomly drawn) pool of knot configurations
2. For each perturbation step. . .
  - 2.1 Randomly move, add, or delete a knot from each set
  - 2.2 If the perturbed knot has a better likelihood  $p(y|k, \xi)$  than the original, replace the original with it
3. For each epoch. . .

---

## Possible evolutionary splines implementation

1. Start with a set of initial (randomly drawn) pool of knot configurations
2. For each perturbation step. . .
  - 2.1 Randomly move, add, or delete a knot from each set
  - 2.2 If the perturbed knot has a better likelihood  $p(y|k, \xi)$  than the original, replace the original with it
3. For each epoch. . .
  - 3.1 Take  $i$  perturbation steps on each model in the pool

---

## Possible evolutionary splines implementation

1. Start with a set of initial (randomly drawn) pool of knot configurations
2. For each perturbation step. . .
  - 2.1 Randomly move, add, or delete a knot from each set
  - 2.2 If the perturbed knot has a better likelihood  $p(y|k, \xi)$  than the original, replace the original with it
3. For each epoch. . .
  - 3.1 Take  $i$  perturbation steps on each model in the pool
  - 3.2 Store a pool of the  $j$  best models

---

## Possible evolutionary splines implementation

1. Start with a set of initial (randomly drawn) pool of knot configurations
2. For each perturbation step. . .
  - 2.1 Randomly move, add, or delete a knot from each set
  - 2.2 If the perturbed knot has a better likelihood  $p(y|k, \xi)$  than the original, replace the original with it
3. For each epoch. . .
  - 3.1 Take  $i$  perturbation steps on each model in the pool
  - 3.2 Store a pool of the  $j$  best models
  - 3.3 Create a new pool by randomly recombining existing pool

---

## Possible evolutionary splines implementation

1. Start with a set of initial (randomly drawn) pool of knot configurations
2. For each perturbation step. . .
  - 2.1 Randomly move, add, or delete a knot from each set
  - 2.2 If the perturbed knot has a better likelihood  $p(y|k, \xi)$  than the original, replace the original with it
3. For each epoch. . .
  - 3.1 Take  $i$  perturbation steps on each model in the pool
  - 3.2 Store a pool of the  $j$  best models
  - 3.3 Create a new pool by randomly recombining existing pool
4. Pool will [*hopefully*] converge to MLE

---

## Possible evolutionary splines implementation

1. Start with a set of initial (randomly drawn) pool of knot configurations
2. For each perturbation step. . .
  - 2.1 Randomly move, add, or delete a knot from each set
  - 2.2 If the perturbed knot has a better likelihood  $p(y|k, \xi)$  than the original, replace the original with it
3. For each epoch. . .
  - 3.1 Take  $i$  perturbation steps on each model in the pool
  - 3.2 Store a pool of the  $j$  best models
  - 3.3 Create a new pool by randomly recombining existing pool
4. Pool will [*hopefully*] converge to MLE
5. Can maximise posterior, instead of likelihood, by including a prior term

---

## BARS vs evolution

- ✶ Would be interesting to compare this to the MCMC method
  - ▶ MCMC sets up a Markov chain whose stationary distribution is the posterior
  - ▶ This aims to find a Markov chain whose stationary distribution is the  $\text{argmax}$
  - ▶ Grounds for a rigorous justification / proof of convergence
- ✶ Evolution will likely be faster
  - ▶ Could leverage existing genetic optimization packages
  - ▶ Easily parallelised for more speed-up
  - ▶ No RJ-MCMC makes it easier target for probabilistic programming

---

## SV-BARS

Another idea: remodel BARS to work similarly to sparse GPR

- ✂ BARS typically uses 3000+ MCMC steps
  - ▶ Each MCMC step requires inverting an  $n \times n$  matrix – SLOW
- ✂ Choose a *[small]* set of maximally informative surrogate datapoints  $(x_i^*, y_i^*)$
- ✂ Run MCMC step on surrogate datapoints
  - ▶ Much faster to invert the smaller matrix
- ✂ Well-chosen surrogate points means we get the same result as running on real data
- ✂ Fewer datapoints means it runs a lot faster



---

## SV-BARS implementation

- Find a set of  $m$  inducing points  $(x_i^*, y_i^*)$ 
  - Find inducing points by minimising  $D_{\text{KL}}[p_y || p_{y^*}]$
  - Use variational Bayes to approximate this
- Find posterior knots  $p(k, \xi | y^*)$  [instead of  $p(k, \xi | y)$ ]
  - $\mathcal{O}(m^3)$ ,  $m \ll n$
  - Sparse GPR is  $\mathcal{O}(nm^3)$ , so if my complexity is correct, we get a bigger speed-up / outperform SVGPR!

Would require learning RJ-MCMC, variational Bayes, sparse GPR in-depth

---

## Periodic BARS

An approach for if we need to model more than a few spikes:

✿ Assume data are given by  $y_i = f(x_i) + \varepsilon$

▶  $f(t) = f(t + T)$

✿ Find either

▶  $T$ -periodic knot-set

▶  $T$ -periodic basis splines

▶ Nicer approach


... in such a way that we ...

✿ minimise fitting time

✿ balance fit against number of knots

---

## Semiperiodic methods




Experimental data aren't perfectly periodic

---

---

## Semiperiodic methods



---

Experimental data aren't perfectly periodic *[could  $\text{Ca}^{2+}$ , or experiment setup!]*

---

---

## Semiperiodic methods

✎ Assume data are given by  $y_i = A(x_i)f(x_i) + \varepsilon$

---

## Semiperiodic methods

✎ Assume data are given by  $y_i = A(x_i)f(x_i) + \varepsilon$

▶  $f(t) = f(t + T)$  is the periodic behaviour

---

## Semiperiodic methods

✿ Assume data are given by  $y_i = A(x_i)f(x_i) + \varepsilon$

- ▶  $f(t) = f(t + T)$  is the periodic behaviour
- ▶  $A(t)$  is the drifting amplitude

---

## Semiperiodic methods

- ✿ Assume data are given by  $y_i = A(x_i)f(x_i) + \varepsilon$
- ▶  $f(t) = f(t + T)$  is the periodic behaviour
  - ▶  $A(t)$  is the drifting amplitude
  - ▶ Might require transforming the data, to a zero-DC offset



---

## Semiperiodic methods

- ✿ Assume data are given by  $y_i = A(x_i)f(x_i) + \varepsilon$ 
  - ▶  $f(t) = f(t + T)$  is the periodic behaviour
  - ▶  $A(t)$  is the drifting amplitude
  - ▶ Might require transforming the data, to a zero-DC offset
- ✿ Fit a model to  $f(t)$

---

## Semiperiodic methods

- ✿ Assume data are given by  $y_i = A(x_i)f(x_i) + \varepsilon$ 
  - ▶  $f(t) = f(t + T)$  is the periodic behaviour
  - ▶  $A(t)$  is the drifting amplitude
  - ▶ Might require transforming the data, to a zero-DC offset
- ✿ Fit a model to  $f(t)$ 
  - ▶ Only requires one period's worth of data

---

## Semiperiodic methods

- ✿ Assume data are given by  $y_i = A(x_i)f(x_i) + \varepsilon$ 
  - ▶  $f(t) = f(t + T)$  is the periodic behaviour
  - ▶  $A(t)$  is the drifting amplitude
  - ▶ Might require transforming the data, to a zero-DC offset
- ✿ Fit a model to  $f(t)$ 
  - ▶ Only requires one period's worth of data
- ✿ Fit a model to  $A(t)$

---

## Semiperiodic methods

- ✿ Assume data are given by  $y_i = A(x_i)f(x_i) + \varepsilon$ 
  - ▶  $f(t) = f(t + T)$  is the periodic behaviour
  - ▶  $A(t)$  is the drifting amplitude
  - ▶ Might require transforming the data, to a zero-DC offset
- ✿ Fit a model to  $f(t)$ 
  - ▶ Only requires one period's worth of data
- ✿ Fit a model to  $A(t)$ 
  - ▶ Requires a few datapoints across all periods

---

## Semiperiodic methods

- ✿ Assume data are given by  $y_i = A(x_i)f(x_i) + \varepsilon$ 
  - ▶  $f(t) = f(t + T)$  is the periodic behaviour
  - ▶  $A(t)$  is the drifting amplitude
  - ▶ Might require transforming the data, to a zero-DC offset
- ✿ Fit a model to  $f(t)$ 
  - ▶ Only requires one period's worth of data
- ✿ Fit a model to  $A(t)$ 
  - ▶ Requires a few datapoints across all periods
- ✿ BARS, GPR, NARMAX all interesting model options

---

## Semiperiodic methods

- ✿ Assume data are given by  $y_i = A(x_i)f(x_i) + \varepsilon$ 
  - ▶  $f(t) = f(t + T)$  is the periodic behaviour
  - ▶  $A(t)$  is the drifting amplitude
  - ▶ Might require transforming the data, to a zero-DC offset
- ✿ Fit a model to  $f(t)$ 
  - ▶ Only requires one period's worth of data
- ✿ Fit a model to  $A(t)$ 
  - ▶ Requires a few datapoints across all periods
- ✿ BARS, GPR, NARMAX all interesting model options
- ✿ Would likely give similar results to sparse BARS

---

## Conclusions

- ✶ For small numbers of cycles (1, 2, 3), current BARS works well enough for CBC

---

## Conclusions

- ✦ For small numbers of cycles (1, 2, 3), current BARS works well enough for CBC
- ✦ For bigger data, we need something more creative



---

## Conclusions

- ✦ For small numbers of cycles (1, 2, 3), current BARS works well enough for CBC
- ✦ For bigger data, we need something more creative
- ✦ Sparse variational BARS would be valuable within machine learning community

---

## Conclusions

- ✦ For small numbers of cycles (1, 2, 3), current BARS works well enough for CBC
- ✦ For bigger data, we need something more creative
- ✦ Sparse variational BARS would be valuable within machine learning community
  - ▶ Would very likely give state-of-the-art results!

---

## Conclusions

- ✦ For small numbers of cycles (1, 2, 3), current BARS works well enough for CBC
- ✦ For bigger data, we need something more creative
- ✦ Sparse variational BARS would be valuable within machine learning community
  - ▶ Would very likely give state-of-the-art results!
  - ▶ Could be combined with the evolutionary approach for more speedup

---

## Conclusions

- ✂ For small numbers of cycles (1, 2, 3), current BARS works well enough for CBC
- ✂ For bigger data, we need something more creative
- ✂ Sparse variational BARS would be valuable within machine learning community
  - ▶ Would very likely give state-of-the-art results!
  - ▶ Could be combined with the evolutionary approach for more speedup
- ✂ (Semi)periodic BARS would be less generally applicable, but potentially faster when applicable

---

## Possible routes

1. Use current BARS setup for a CBC experiment

---

## Possible routes

1. Use current BARS setup for a CBC experiment
  - ▶ Gets CBC results fast

## Possible routes

1. Use current BARS setup for a CBC experiment
  - ▶ Gets CBC results fast
2. Adapt current BARS setup for sparsity / evolution, then use in CBC

---

## Possible routes

1. Use current BARS setup for a CBC experiment
  - ▶ Gets CBC results fast
2. Adapt current BARS setup for sparsity / evolution, then use in CBC
  - ▶ New splines method would be valuable in ML community



---

## Possible routes

1. Use current BARS setup for a CBC experiment
  - ▶ Gets CBC results fast
2. Adapt current BARS setup for sparsity / evolution, then use in CBC
  - ▶ New splines method would be valuable in ML community
3. Demonstrate splines [*and GPR?*] on other problems (NDC, comp-synth-bio, ML)

---

## Possible routes

1. Use current BARS setup for a CBC experiment
  - ▶ Gets CBC results fast
2. Adapt current BARS setup for sparsity / evolution, then use in CBC
  - ▶ New splines method would be valuable in ML community
3. Demonstrate splines [*and GPR?*] on other problems (NDC, comp-synth-bio, ML)
  - ▶ Extra paper, builds well on the surrogate-modelling knowledge I'm developing

---

## Possible routes

1. Use current BARS setup for a CBC experiment
  - ▶ Gets CBC results fast
2. Adapt current BARS setup for sparsity / evolution, then use in CBC
  - ▶ New splines method would be valuable in ML community
3. Demonstrate splines [*and GPR?*] on other problems (NDC, comp-synth-bio, ML)
  - ▶ Extra paper, builds well on the surrogate-modelling knowledge I'm developing
  - ▶ I wouldn't know which problems to apply them to

---

## Possible routes

1. Use current BARS setup for a CBC experiment
  - ▶ Gets CBC results fast
2. Adapt current BARS setup for sparsity / evolution, then use in CBC
  - ▶ New splines method would be valuable in ML community
3. Demonstrate splines [*and GPR?*] on other problems (NDC, comp-synth-bio, ML)
  - ▶ Extra paper, builds well on the surrogate-modelling knowledge I'm developing
  - ▶ I wouldn't know which problems to apply them to
4. Make a periodic BARS setup, then use that in CBC

---

## Possible routes

1. Use current BARS setup for a CBC experiment
  - ▶ Gets CBC results fast
2. Adapt current BARS setup for sparsity / evolution, then use in CBC
  - ▶ New splines method would be valuable in ML community
3. Demonstrate splines [*and GPR?*] on other problems (NDC, comp-synth-bio, ML)
  - ▶ Extra paper, builds well on the surrogate-modelling knowledge I'm developing
  - ▶ I wouldn't know which problems to apply them to
4. Make a periodic BARS setup, then use that in CBC
  - ▶ Periodic BARS is a less general method

---

## My proposal

- ✂ Validate models
  - ▶ Also try simple data transformations for GPR
- ✂ Write up notes on *everything* so far

Then...

1. Adapt BARS for sparsity, evolution
  - ▶ Fast, SOTA, scalable
  - ▶ Needs variational Bayes learned
2. Demonstrate sparse BARS on other problems (NDC, comp-synth-bio, ML)
  - ▶ Not sure which problems would benefit from splines?
3. Apply the shiny new splines method to CBC